



THE PATTERNS OF DISTRIBUTED LOGGING AND CONTAINERS

CloudNativeCon Europe 2017
March 30, 2017

Satoshi Tagomori (@tagomoris)
Treasure Data, Inc.



TREASURE DATA



SATOSHI TAGOMORI

(@tagomoris)

Fluentd, MessagePack-Ruby, Norikra, ...

Treasure Data, Inc.



**TREASURE
DATA**



CLOUD NATIVE
COMPUTING FOUNDATION

1. Microservices, Containers and Logging
2. Scaling Logging Platform
3. Patterns: Source/Destination -side Aggregation
4. Patterns: Scaling Up/Out Destination
5. Practices

MICROSERVICES, CONTAINERS AND LOGGING

Logging in Industries

- Service Logs
 - Web access logs
 - Ad logs
 - Commercial transaction logs for analytics (EC, Game, ...)
- System Logs
 - Syslog and other OS logs
 - Audit logs
 - Performance metrics

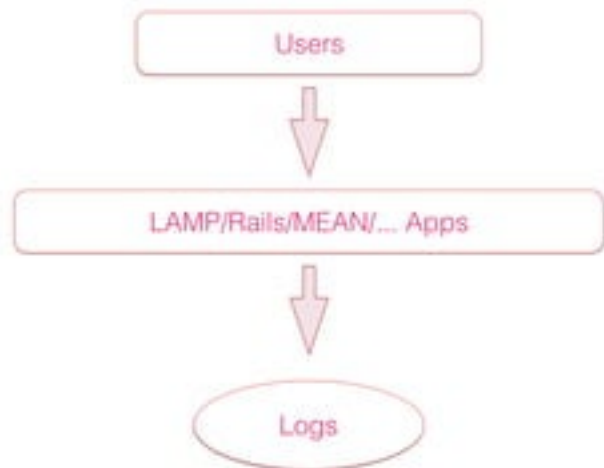


Logs for
Business Growth



Logs for
Service Stability

Microservices and Logging



Monolithic service



Microservices

Microservices and Containers

- Microservices
 - Isolated dependencies
 - Agile deployment
- Containers
 - Isolated environments & resources
 - Simple pull&restart deployment
 - Less overhead, high density

Logging Challenges with Microservices/Containers

- Containerization changes everything:
 - No permanent storages
 - No fixed physical/network addresses
 - No fixed mapping between servers and roles

Logging Challenges with Microservices/Containers

- Containerization changes everything:
 - No permanent storages
 - No fixed physical/network addresses
 - No fixed mapping between servers and roles

Transfer Logs to Anywhere ASAP

Logging Challenges with Microservices/Containers

- Containerization changes everything:
 - No permanent storages
 - No fixed physical/network addresses
 - No fixed mapping between servers and roles

Push Logs From Containers

Logging Challenges with Microservices/Containers

- Containerization changes everything:
 - No permanent storages
 - No fixed physical/network addresses
 - No fixed mapping between servers and roles

Label Logs With Service Names/Tags

Logging Challenges with Microservices/Containers

- Containerization changes everything:
 - No permanent storages
 - No fixed physical/network addresses
 - No fixed mapping between servers and roles

Label Logs With Service Names/Tags

Parse Logs & Label Values At Source

Structured Logs

Structured Logs: tag, time, key-value pairs

Original log:

the customer put an item to cart: item_id=101, items=10, client=web

Structured log:

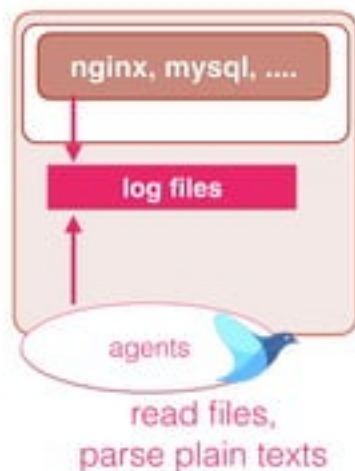
ec_service.shopping_cart tag

2017-03-30 16:35:37 +0100 timestamp

```
{
  "container_id": "bfdd5b9....",
  "container_name": "/infallible_mayer",
  "source": "stdout",
  "event": "put an item to cart",
  "item_id": 101,
  "items": 10,
  "client": "web"
}
```

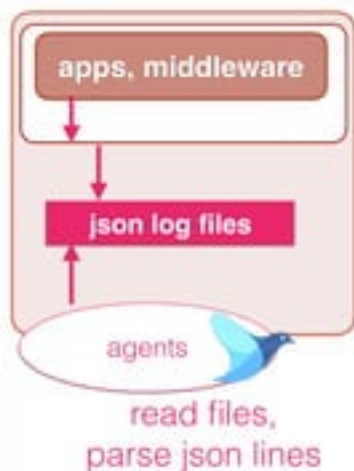
) record

How to Ship Logs from Docker Containers



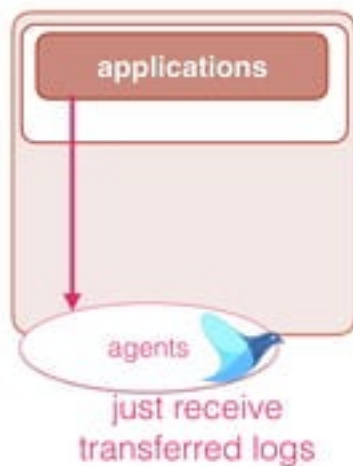
**Using
mounted volume**

+ disk I/O penalty
+ mount points




**Using
container json logs**

+ disk I/O penalty



**Sending logs
to agents directly**

+ logger code
+ agent config




**Using
logging drivers**





SCALING LOGGING PLATFORM

Core Architecture: Distributed Logging



Distributed Logging Workflow

Collector

- Retrieve raw logs: file system / network
- Parse log content

Aggregator

- Get data from multiple sources
- Split/merge incoming data into streams

Destination

- Retrieve structured logs from Aggregator
- Store formatted logs

Scaling Logging

- Network Traffic
 - Split heavy log traffic into traffics to nodes
- CPU Load
 - Distribute processing to nodes about parsing/formatting logs
- High Availability
 - Switch traffic from a node to another for failures
- Agility
 - Reconfigure whole logging layer to modify destinations



PATTERNS: SOURCE/DESTINATION -SIDE AGGREGATION

Source Side Aggregation

NO



YES

NO



Destination
Side
Aggregation

YES




Now I'm Talking About:




Source-side Aggregation Patterns

Without Source-side Aggregation



With Source-side Aggregation




Aggregation Pattern without Source-side Aggregation

- **Pros:**

- Simple configuration

- **Cons:**

- Fixed aggregator (destination endpoint) address configured in containers
- Many network connections
- High load in aggregator / destination




Aggregation Pattern with Source-side Aggregation

- **Pros:**

- Less connections
- Lower load in aggregator / destination
- Less configurations in containers
- More agility
(aggregate containers can be reconfigured)


- **Cons:**

- Need more resources (+1 container per host)




Destination-side Aggregation Patterns

**Without
Destination-side Aggregation**



**With
Destination-side Aggregation**




Aggregation Pattern without Destination-side Aggregation

- **Pros:**

- Less nodes
- Simpler configuration

- **Cons:**

- Destination changes affects all source nodes
- Worse performance:
many small write requests on destination(storage)




Aggregation Pattern with Destination-side Aggregation

- **Pros:**

- Destination changes does NOT affect source nodes
- Better performance:
destination aggregator can merge write operations


- **Cons:**

- More nodes
- More complex configuration




PATTERNS: SCALING UP/OUT DESTINATION

Now I'm Talking About:




Scaling Destination Patterns

Scaling Up Aggregator/Destination Endpoints



**Using HTTP Load Balancer
or Huge Queues**

Scaling Out Aggregator/Destination Endpoints



Using Round Robin Clients


Scaling Up Destination

- **Pros:**

- Simple configuration:
specifying load balancer only
in collector nodes

- **Cons:**

- Upper limits about scaling up
on Load balancer (or queue)




Scaling Out Destination

- **Pros:**

- Unlimited scaling by adding nodes

- **Cons:**

- Complex configuration in collector nodes
- Client feature required for round-robin
- Unavailable for traffic over Internet



Destination-side Aggregation and Destination Scaling

Destination Side Aggregation

NO

YES

**Scaling Up
Destination
Endpoints**

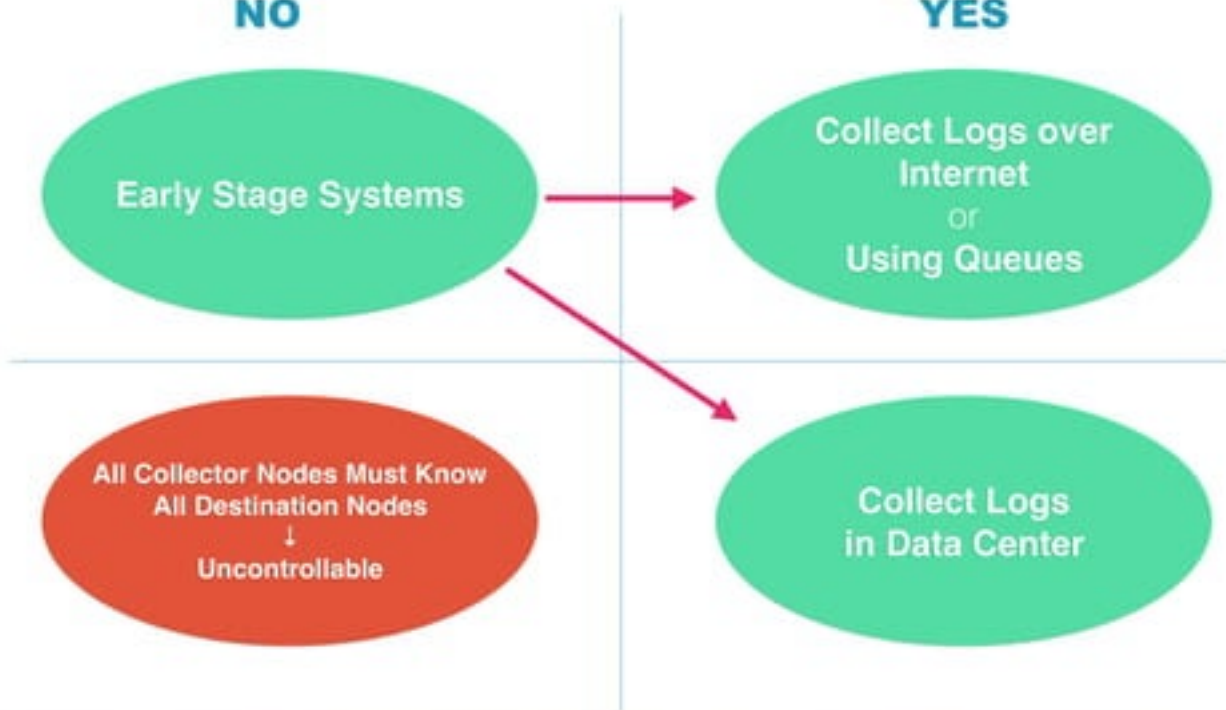
Early Stage Systems

Collect Logs over
Internet
or
Using Queues

**Scaling Out
Destination
Endpoints**

All Collector Nodes Must Know
All Destination Nodes
↓
Uncontrollable

Collect Logs
in Data Center



PRACTICES

Practices: Docker + Fluentd



- **Docker Fluentd Logging Driver**

- Docker containers can send these logs to Fluentd directly, with less overhead

- Fluentd's **Pluggable Architecture**

- Various destination systems (storage/database/service) are available by changing configuration

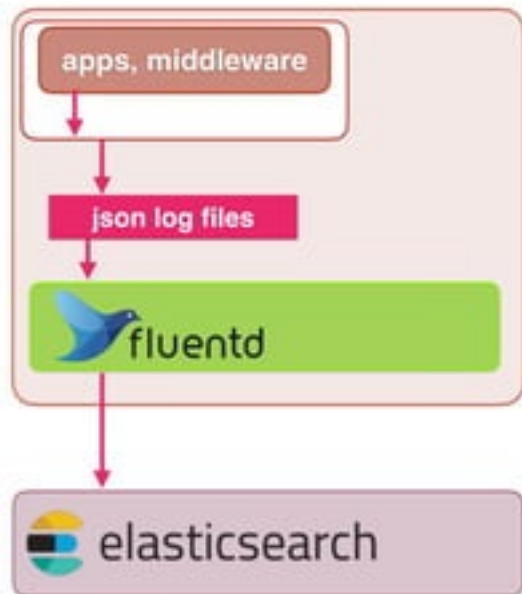
- **Small Memory Footprint**

- Source aggregation requires +1 container per hosts: less additional resource usage is fine!

Practice 1: Source-side Aggregation + Scaling Up

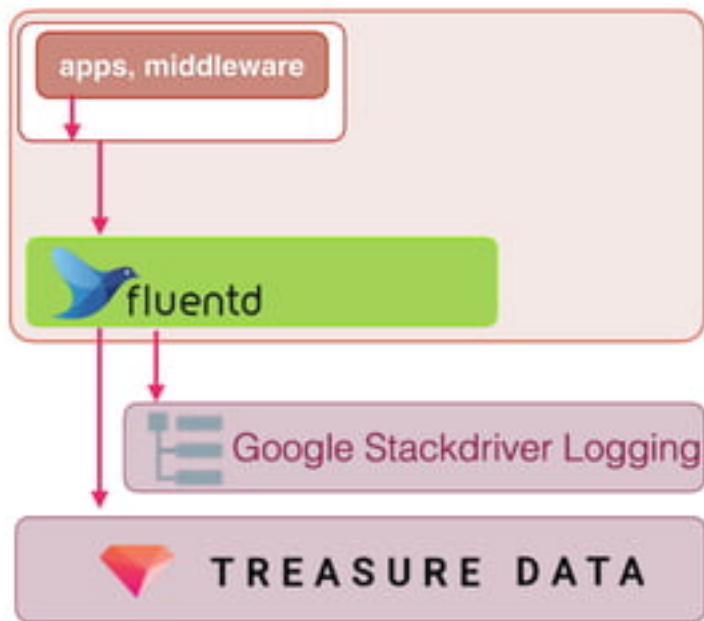


- Kubernetes: Fluentd + Elasticsearch
 - a.k.a EFK stack (inspired by ELK stack)
 - Elasticsearch - Fluentd - Kibana



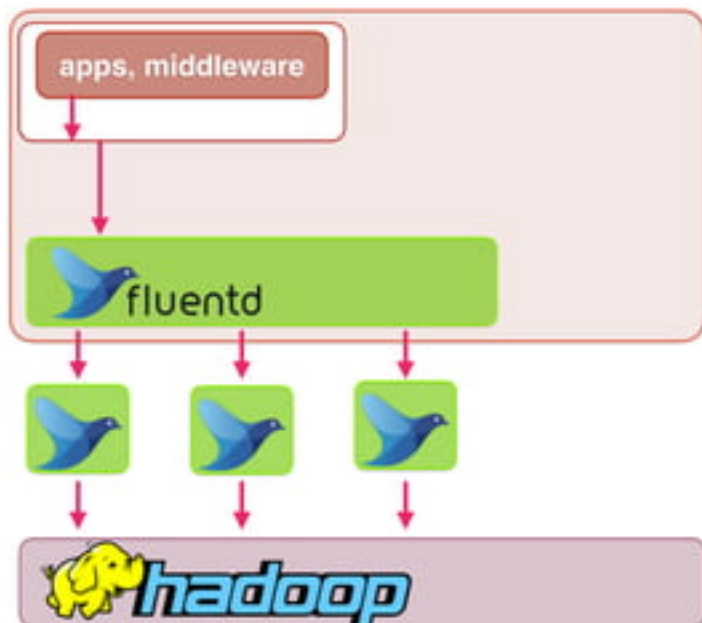
Practice 2: Source-side Aggregation + Scaling Up

- Containerized Applications
 - w/ Google Stackdriver for Monitoring
 - w/ Treasure Data for Analytics



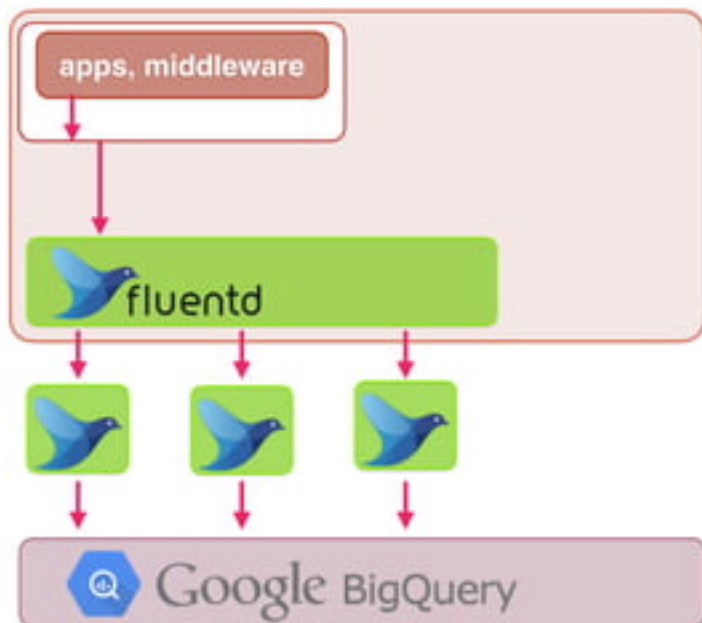
Practice 3: Source/Destination-side Aggregation + Scaling Out

- Containerized Application
 - w/ Log processing on Hadoop
 - writing files on HDFS via WebHDFS
- Hadoop HDFS prefers large files on HDFS:
 - Destination-side aggregation works well



Practice 4: Source/Destination-side Aggregation + Scaling Out

- Containerized Application
 - w/ Log processing on Google BigQuery
 - putting logs via HTTPS
- BigQuery has quota about write requests:
 - Destination-side aggregation works well



Best practices?

- Source aggregation: do it
 - it makes app containers free from logging problems (buffering, HA, ...)
- Destination aggregation: it depends
 - no need for cloud logging services/storages
 - may need for self-hosted distributed filesystems/databases
 - may need for cloud services which charges per requests
- Destination scaling: it depends on destinations

A stylized bird logo composed of overlapping geometric shapes in various shades of blue and cyan. The bird is facing right, with a small white circle for an eye. The background is a solid dark blue.

Make Logging Scalable,
Service Stable & Business Growing.

Happy Logging!
[@tagomoris](#)